

Bayesian Hierarchical MPT Models

Applications with TreeBUGS

Daniel W. Heck



25.02.2020

TreeBUGS: Bayesian Hierarchical MPT Modeling

- 1) The R package TreeBUGS
- 2) Basic modeling
 - Model fitting
 - Convergence
 - Plots
 - Model fit

Software for Hierarchical MPTs

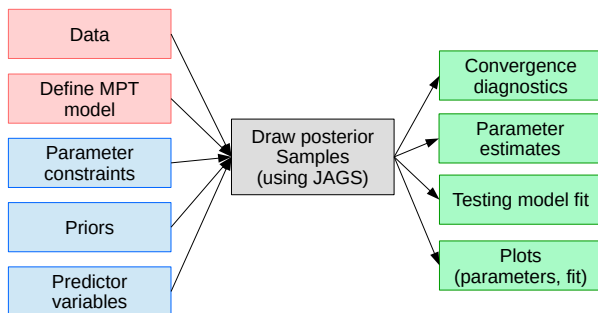
- Implementation of MCMC sampling in R/C/Fortran (Klauer 2010)
- General-purpose software: WinBUGS/JAGS/Stan (Matzke et al. 2015)
- Requires re-implementation of summaries, statistics, plots

TreeBUGS: A user-friendly R package (Heck, Arnold, and Arnold 2018)

- Easy-to-use, open source, free
- Fitting and testing MPT models
 - Posterior sampling, summary statistics, and plots
 - Data generation, robustness simulations
 - Change priors, add predictors, etc.
 - Current limitation: Crossed random effects (for persons & items)

Functionality of TreeBUGS

- Input: R objects or text/csv files (minimal R knowledge required)
- Priors and other details can be changed in R
- TreeBUGS translates the model to JAGS (Plummer, 2003) to draw posterior samples
- Functions for post-processing, summaries and plots



TreeBUGS Paper

Behav Res

DOI 10.3758/s13428-017-0869-7



TreeBUGS: An R package for hierarchical multinomial-processing-tree modeling

Daniel W. Heck¹ · Nina R. Arnold¹ · Denis Arnold^{2,3}

© The Author(s) 2017. This article is published with open access at Springerlink.com

Abstract Multinomial processing tree (MPT) models are a class of measurement models that account for categorical data by assuming a finite number of underlying cognitive processes. Traditionally, data are generated across participants and

estimates, fit statistics, and within- and between-subjects comparisons, as well as goodness-of-fit and summary plots. We also propose and implement novel statistical extensions to include continuous and discrete predictors (e.g., either fixed or

Basic Modeling

(corresponding R script: 04-application-TreeBUGS.R)

Modeling with TreeBUGS is simple

- 1 Specify model and data
- 2 Draw MCMC samples
- 3 Check convergence
- 4 Check model fit
- 5 Interpret/plot parameters

Note that these are the usual steps in any Bayesian analysis...


MPT Model Specification

- MPT structure is defined in an EQN model file
 - Can be copied from multiTree
- Difference: the symbol # allows to add comments

Two-high threshold model (file: 2htm.eqn)

```
# Targets
target hit do
target hit (1-do)*g
target miss (1-do)*(1-g)

# Lures
lure cr dn
lure fa (1-dn)*g
lure cr (1-dn)*(1-g)
```



Data Structure

- Data: Response frequencies in wide format
 - One line per person
 - One category per column
 - Column names must be identical to the EQN categories!
- Either supplied in .csv-file or as `data.frame` / `matrix` in R

Example: 2htm.csv


```
# set working directory:  
# setwd("D:/R/MPT-workshop/")  
frequencies <- read.csv("2htm.csv")  
head(frequencies, 5)
```

```
##    cr fa hit miss  
## 1 47  3  45    5  
## 2 38 12  22   28  
## 3 37 13  21   29  
## 4 37 13  32   18  
## 5 46  4  44    6
```


Heterogeneity

- Load TreeBUGS and plot heterogeneity
- If response frequencies are homogeneous, standard (fixed-effects) MPT models are statistically more efficient

```
library(TreeBUGS)  
plotFreq(frequencies, eqnfile = "2htm.eqn")
```



```
plotFreq(frequencies, boxplot = FALSE, eqnfile = "2htm.eqn")
```



Fitting MPT Models

■ Fitting an MPT model in TreeBUGS

- Model: Text file in EQN syntax (with model equations)
- Data: .csvfile
- Constraints: text file with equality constraints

```
fit <- traitMPT(eqnfile = "htm.txt",  
               data = "responses.csv",  
               restrictions = "2htm_constraints.txt")  
  
# beta-MPT: different function, but identical arguments  
fit_beta <- betaMPT(eqnfile = "htm.txt",  
                   data = "responses.csv",  
                   restrictions = "2htm_constraints.txt")
```

Fitting an MPT Model in R

- Alternative: define everything directly in R
 - Model: Text string (character in apostrophes)
 - Data: Matrix or data frame
 - Constraints: A list

```
htm <- "  
target hit do  
target hit (1-do)*g  
target miss (1-do)* (1-g)  
  
lure cr dn  
lure fa (1-dn)*g  
lure cr (1-dn)*(1-g)  
"  
fit <- traitMPT(eqnfile = htm,  
                data = frequencies,  
                restrictions = list("dn=do", "g=.50"))
```

Parameter Constraints

Equality constraints

```
# (A) use a general model file and constrain parameters:
fit <- traitMPT(eqnfile = htm,
               data = frequencies,
               restrictions = list("dn=do", "g=.50"))





# (B) hard-coding of constraints in the EQN file:
htm_constr <- "
target hit d
target hit (1-d)*.50
target miss (1-d)*.50

lure cr d
lure fa (1-d)*.50
lure cr (1-d)*.50
"
fit <- traitMPT(eqnfile = htm_constr,
               data = frequencies)
```

Convergence

- Convergence check
 - Posterior/ MCMC samples should look unsystematic (like a hairy caterpillar)
 - For more options, see: `?plot.traitMPT`


```
plot(fit, parameter = "mean", type = "default")
```




Convergence

■ Interpreting MCMC plots

Good convergence: 'hairy caterpillar'




Bad convergence: slow movement




Convergence

- Gelman-Rubin statistic (\hat{R})
 - Also known as: “potential scale reduction factor” or “R hat”
 - Similar to ANOVA: Compares between-chain and within-chain variances (large differences between these variances indicate nonconvergence)
 - Statistic should be close to 1 (standard criterion: $\hat{R} < 1.05$)

Good convergence: all chains similar



Bad convergence: chains differ



Convergence

- Gelman-Rubin statistic (\hat{R})
 - Columns Rhat and R_95% in the summary output

```
summary(fit)
```

```
## Call:
## traitMPT(eqnfile = htm, data = frequencies, restrictions = list("dn=do"),
##       ppp = 1000)
##
## Group-level medians of MPT parameters (probability scale):
##           Mean    SD  2.5%   50% 97.5% Time-series SE n.eff  Rhat R_95%
## mean_dn 0.536 0.060 0.414 0.537 0.654          0.003   354 1.027 1.094
## mean_g  0.475 0.031 0.415 0.475 0.537          0.001   685 1.005 1.016
##
## Mean/Median of latent-trait values (probit-scale) across individuals:
##           Mean    SD  2.5%   50% 97.5% Time-series SE n.eff  Rhat R_95%
## latent_mu_dn 0.092 0.154 -0.216 0.093 0.396          0.008   354 1.027 1.095
## latent_mu_g -0.062 0.078 -0.215 -0.062 0.093          0.003   684 1.005 1.016
##
## Standard deviation of latent-trait values (probit scale) across individuals:
##           Mean    SD  2.5%   50% 97.5% Time-series SE n.eff  Rhat R_95%
## latent_sigma_dn 1.062 0.137 0.828 1.052 1.368          0.002  3439 1.005 1.019
## latent_sigma_g  0.433 0.070 0.311 0.428 0.585          0.001  2890 1.005 1.018
##
## Correlations of latent-trait values on probit scale:
##           Mean    SD  2.5%   50% 97.5% Time-series SE n.eff  Rhat R_95%
## rho[dn,g] 0.315 0.17 -0.041 0.323 0.617          0.003  2380 1.002 1.005
##
## Correlations (posterior mean estimates) in matrix form:
##           dn      g
```

Options for MCMC Sampling

- If the model has not converged, it must be fitted with more conservative settings:

```
fit <- traitMPT(  
  eqnfile = htm, data = frequencies,  
  restrictions = list("dn=do"),  
  
  n.adapt = 5000, # longer adaption of JAGS increases efficiency of sampling  
  n.burnin = 5000, # longer burnin avoids issues due to bad starting values  
  n.iter = 30000, # drawing more MCMC samples leads to higher precision  
  n.thin = 10,    # ommitting every 10th sample reduces memory load  
  n.chains = 4)   # more MCMC chains increase precision
```

```
## MCMC sampling started at 2020-02-20 12:38:26  
## Calling 4 simulations using the parallel method...  
## Following the progress of chain 1 (the program will wait for all chains  
## to finish before continuing):  
## Welcome to JAGS 4.3.0 on Thu Feb 20 12:38:29 2020  
## JAGS is free software and comes with ABSOLUTELY NO WARRANTY  
## Loading module: basemod: ok  
## Loading module: bugs: ok  
## . Loading module: dic: ok  
## . Loading module: glm: ok  
## . . Reading data file data.txt  
## . Compiling data graph  
## Resolving undeclared variables
```

Extend MCMC Sampling

- If the MCMC samples are OK but higher precision is needed:
 - Extend sampling and add new MCMC samples to the fitted JAGS object





```
fit2 <- extendMPT(fit,           # fitted MPT model
                 n.adapt = 2000, # JAGS need to restart and adapt again
                 n.burnin = 0,   # burnin not needed if previous samples are OK
                 n.iter = 10000) # how many additional iterations?
```

```
## Calling 4 simulations using the parallel method...
## Following the progress of chain 1 (the program will wait for all chains
## to finish before continuing):
## Welcome to JAGS 4.3.0 on Thu Feb 20 12:38:58 2020
## JAGS is free software and comes with ABSOLUTELY NO WARRANTY
## Loading module: basemod: ok
## Loading module: bugs: ok
## . Loading module: dic: ok
## . Loading module: glm: ok
## . . Reading data file data.txt
## . Compiling data graph
##   Resolving undeclared variables
##   Allocating nodes
##   Initializing
##   Reading data back into data table
## Compiling model graph
##   Resolving undeclared variables
##   Allocating nodes
## Graph information:
##   Observed stochastic nodes: 100
##   Unobserved stochastic nodes: 55
##   Total graph nodes: 155
```

Convergence for Second Fit

- Check convergence again:

```
plot(fit2, parameter = "mean", type = "default")
```



Parameter Estimates

- Summary statistics for posterior distribution:
 - Posterior mean and median (50% quantile)
 - Posterior standard deviation (SD, similar to standard error)
 - Bayesian credibility interval (2.5% and 97.5% quantiles)

```
summary(fit2)
```


```
## Call:
## [[1]]
## traitMPT(eqnfile = htm, data = frequencies, restrictions = list("dn=do"),
##   n.iter = 30000, n.adapt = 5000, n.burnin = 5000, n.thin = 10,
##   n.chains = 4)
##
## [[2]]
## extendMPT(fittedModel = fit, n.iter = 10000, n.adapt = 2000,
##   n.burnin = 0)
##
##
## Group-level medians of MPT parameters (probability scale):
##      Mean    SD  2.5%   50% 97.5% Time-series SE n.eff  Rhat R_95%
## mean_dn 0.540 0.063 0.417 0.540 0.668      0.002   698 1.006 1.018
## mean_g  0.474 0.031 0.415 0.474 0.537      0.001  1595 1.002 1.005
##
## Mean/Median of latent-trait values (probit-scale) across individuals:
##      Mean    SD  2.5%   50% 97.5% Time-series SE n.eff  Rhat R_95%
## latent_mu_dn 0.102 0.162 -0.210 0.100 0.435      0.006   696 1.006 1.018
## latent_mu_g -0.064 0.078 -0.216 -0.066 0.094      0.002  1594 1.002 1.005
##
## Standard deviation of latent-trait values (probit scale) across individuals:
##      Mean    SD  2.5%   50% 97.5% Time-series SE n.eff  Rhat R_95%
## latent_sigma_dn 1.066 0.14 0.828 1.054 1.376      0.002  4662 1.004 1.011
## latent_sigma_g  0.432 0.07 0.311 0.428 0.586      0.001  6041 1.000 1.001
```

Plot Parameter Estimates

- Estimates for group-level parameters
 - Overall mean μ : Posterior mean and Bayesian credibility interval
 - Individual parameters θ_i : Posterior mean

```
plotParam(fit)
```

group-level means + 95% CI (red) and individual means




Plot Parameter Estimates

- Plot parameter profiles per person
- For instance, to assess the test-retest reliability of a parameter (Michalkiewicz & Erdfelder, 2016)
- For more options, see: `?plotParam`

```
plotParam(fit, addLines = TRUE, select = c("dn", "g"))
```

roup-level means + 95% CI (red) and individual means




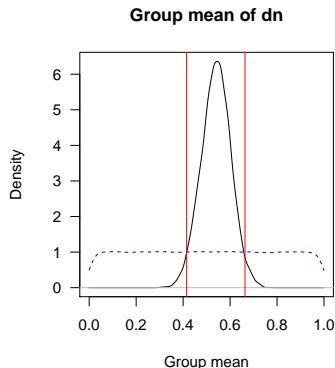
Compare Prior and Posterior

How much did we learn about the parameters?

- Graphical assessment: Plot prior (blue) and posterior (black) densities

```
plotPriorPost(fit)
```

```
## Press <Enter> to show the next plot.
```



```
## Press <Enter> to show the next plot.
```


Model Fit: Predicted vs. Observed Data


Graphical test of model fit

- Plot means of observed frequencies
- Compare against posterior-predicted frequencies (boxplots)

```
colMeans(frequencies)  # observed group means that are tested
```

```
##      cr      fa      hit      miss  
## 39.74 10.26 36.98 13.02
```

```
plotFit(fit)           # graphical test
```



Model Fit: Predicted vs. Observed Data


■ Model fit for *covariances* of the observed frequencies

```
cov(frequencies)           # observed covariance matrix that is tested
```

```
##           cr           fa           hit           miss
## cr      55.38000 -55.38000  41.21918 -41.21918
## fa     -55.38000  55.38000 -41.21918  41.21918
## hit     41.21918 -41.21918 106.02000 -106.02000
## miss   -41.21918  41.21918 -106.02000 106.02000
```

```
plotFit(fit, stat = "cov") # graphical test
```

Observed (red) and predicted (gray) covariances



Model Fit: Test Statistic for Predicted vs. Observed Data

Testing model fit

- We need test statistics to quantify model fit (Klauer, 2010)
 - T1 statistic: Mean structure of frequencies
 - T2 statistic: Covariance matrix of frequencies
 - Similar to Pearson's χ^2 : discrepancy between observed and expected data
- Posterior predictive p -value (PPP) measures model fit:
 - 1) Compute T1 for the observed data
 - 2) Compute T1 for the posterior predicted data
 - 3) PPP = probability that T1(predicted) is larger than T1(observed)
- Interpretation
 - Values around .50 indicate good model fit
 - Values close to 0 indicate misfit
 - In contrast to frequentist p -values, PPP values are *not* uniformly distributed when generating data from the correct model

```
PPP(fit, M = 2000, nCPU = 4)
```

```
## ## Mean structure (T1):
## Observed = 0.0325821 ; Predicted = 0.03440444 ; p-value = 0.513
##
## ## Covariance structure (T2):
## Observed = 7.316926 ; Predicted = 7.525546 ; p-value = 0.51
##
## ## Individual fit (T1):
##      1      2      3      4      5      6      7      8      9     10     11     12     13
## 0.530 0.480 0.508 0.526 0.542 0.496 0.534 0.518 0.504 0.417 0.538 0.335 0.378
##      14     15     16     17     18     19     20     21     22     23     24     25     26
```

Appendix

Appendix: Testing for Heterogeneity

■ Test by Smith and Batchelder (2008)

A) Test person heterogeneity assuming items homogeneity (χ^2)

B) Test person heterogeneity under item heterogeneity (permutation bootstrap)

```
# A) chi^2 test
test <- testHetChi(freq = frequencies,
                  tree = c(cr="lure", fa="lure",
                          hit="target", miss="target"))
data.frame(test)
```

```
##      chisq df      prob
## 1 872.2505 98 1.782941e-124
```


```
# B) requires data in long format (variables: person / item / response)
# testHetPerm(data, tree, source = "person")
```

Appendix: Group-Level Distribution

Distribution of individual estimates

- Histogram: Distribution of θ estimates (posterior mean per person)
- Red density: Estimated group-level distribution

```
plotDistribution(fit) # graphical test
```




Appendix: Convergence

Autocorrelation function

- How strongly are the MCMC samples correlated between iteration t and iteration $t + \text{Lag}$?
- Ideally, these curves should rapidly decrease towards zero.

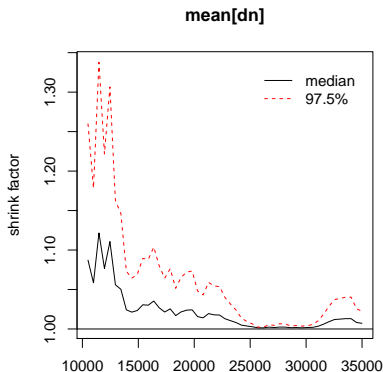
```
plot(fit, parameter = "mean", type = "acf")
```




Appendix: Convergence

- Plot evolution of Gelman-Rubin statistic
 - Also known as: “potential scale reduction factor” or “R hat”
 - Similar to ANOVA: Compares between-chain and within-chain variances (large differences between these variances indicate nonconvergence)
 - Statistic should be close to 1

```
plot(fit, parameter = "mean", type = "gelman")
```



last iteration in chain



last iteration in chain

MCMC samplers available in TreeBUGS

- Fixed-effects (“standard”) MPT: Gibbs sampler in C++
- Beta-MPT: JAGS and C++
- Latent-trait with extensions: JAGS
 - Combination of random- and fixed-effects parameters
 - Continuous and categorical covariates
 - Group-level structure: Independent normal distributions

References I

Heck, Daniel W, Nina R. Arnold, and Denis Arnold. 2018. "TreeBUGS: An R Package for Hierarchical Multinomial-Processing-Tree Modeling." *Behavior Research Methods* 50: 264–84. <https://doi.org/10.3758/s13428-017-0869-7>.

Klauer, K. C. 2010. "Hierarchical Multinomial Processing Tree Models: A Latent-Trait Approach." *Psychometrika* 75: 70–98.
<https://doi.org/10.1007/s11336-009-9141-0>.

Matzke, Dora, Conor V. Dolan, William H. Batchelder, and Eric-Jan Wagenmakers. 2015. "Bayesian Estimation of Multinomial Processing Tree Models with Heterogeneity in Participants and Items." *Psychometrika* 80: 205–35. <https://doi.org/10.1007/s11336-013-9374-9>.

Smith, J. B., and W. H. Batchelder. 2008. "Assessing Individual Differences in Categorical Data." *Psychonomic Bulletin & Review* 15: 713–31.
<https://doi.org/10.3758/PBR.15.4.713>.