

# Bayesian Hierarchical MPT Modeling

## Application and Practice with TreeBUGS

Daniel W. Heck



2023-04-23

# TreeBUGS: Bayesian Hierarchical MPT Modeling

- 1) The R package TreeBUGS
- 2) Basic modeling
  - Model fitting
  - Convergence
  - Plots
  - Model fit
- 3) Advanced modeling
  - Within-subject comparisons
  - Between-subject comparisons
  - Covariates
- 4) Sensitivity/robustness analysis
  - Priors
  - Predictive distributions
  - Simulation

## Software for Hierarchical MPTs

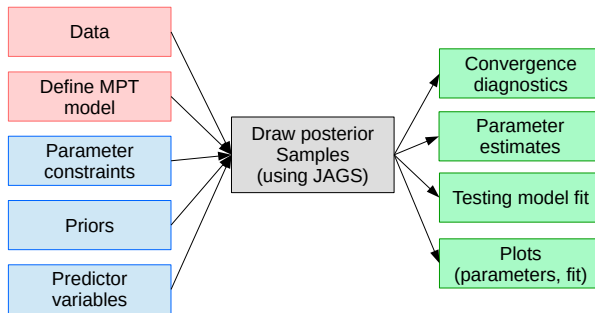
- Implementation of MCMC sampling in R/C/Fortran (Klauer 2010)
- General-purpose software: WinBUGS/JAGS/Stan (Matzke et al. 2015)
- Requires re-implementation of summaries, statistics, plots

## **TreeBUGS: A user-friendly R package** (Heck, Arnold, and Arnold 2018)

- Easy-to-use, open source, free
- Fitting and testing MPT models
  - Posterior sampling, summary statistics, and plots
  - Data generation, robustness simulations
  - Change priors, add predictors, etc.
  - Current limitation: Crossed random effects (persons + items)

## Functionality of TreeBUGS

- Input: R objects or text/csv files (minimal R knowledge required)
- Priors and other details can be changed in R
- TreeBUGS translates the model to JAGS (Plummer, 2003) to draw posterior samples
- Functions for post-processing, summaries and plots



## TreeBUGS Paper

Behav Res

DOI 10.3758/s13428-017-0869-7



### TreeBUGS: An R package for hierarchical multinomial-processing-tree modeling

Daniel W. Heck<sup>1</sup> · Nina R. Arnold<sup>1</sup> · Denis Arnold<sup>2,3</sup>

© The Author(s) 2017. This article is published with open access at Springerlink.com

**Abstract** Multinomial processing tree (MPT) models are a class of measurement models that account for categorical data by assuming a finite number of underlying cognitive processes. Traditionally, data are aggregated across participants and

estimates, fit statistics, and within- and between-subjects comparisons, as well as goodness-of-fit and summary plots. We also propose and implement novel statistical extensions to include continuous and discrete predictors (as either fixed or

## Basic Modeling

(corresponding R script: 06-ApplicationIII-TreeBUGS.R)

# MPT Model Specification

- MPT structure is defined in an EQN model file
  - Simple: copy from multiTree :-)
- Difference: the symbol # allows to add comments


## Two-high threshold model (file: 2htm.eqn)

### # Targets

```
target hit do
target hit (1-do)*g
target miss (1-do)* (1-g)
```

### # Lures

```
lure cr dn
lure fa (1-dn)*g
lure cr (1-dn)*(1-g)
```



- Data: Response frequencies in wide format
  - either .csv-file or data.frame / matrix in R
  - One line per person
  - One category per column
  - Column names must be identical to the EQN categories!

## Example: 2htm.csv

```
# set working directory:  
# setwd("D:/R/MPT-workshop/")  
frequencies <- read.csv("2htm.csv")  
head(frequencies, 5)
```


```
##    cr fa hit miss  
## 1 38 12 36  14  
## 2 33 17 41   9  
## 3 50  0 50   0  
## 4 46  4 36  14  
## 5 29 21 45   5
```




# Heterogeneity

## ■ Load TreeBUGS and plot heterogeneity

```
library("TreeBUGS")  
plotFreq(frequencies, eqnfile = "2htm.eqn")
```



```
plotFreq(frequencies, boxplot = FALSE, eqnfile = "2htm.eqn")
```



# Fitting MPT Models

- Fitting an MPT model in TreeBUGS
  - Model: Text file in EQN syntax (with model equations)
  - Data: .csv file
  - Constraints: text file with equality constraints

```
fit <- traitMPT(eqnfile = "htm.txt",  
               data = "responses.csv",  
               restrictions = "2htm_constraints.txt")  
  
# beta-MPT: different function, but identical arguments  
fit_beta <- betaMPT(eqnfile = "htm.txt",  
                   data = "responses.csv",  
                   restrictions = "2htm_constraints.txt")
```

# Fitting an MPT Model in R

- Alternative: define everything directly in R
  - Model: Text string (character in apostrophes)
  - Data: Matrix or data frame
  - Constraints: A list

```
htm <- "  
target hit do  
target hit (1-do)*g  
target miss (1-do)* (1-g)  
  
lure cr dn  
lure fa (1-dn)*g  
lure cr (1-dn)*(1-g)  
"  
fit <- traitMPT(eqnfile = htm,  
               data = frequencies,  
               restrictions = list("dn=do", "g=.50"))
```

# Parameter Constraints

## Equality constraints

*# (A) use a general model file and constrain parameters:*

```
fit <- traitMPT(eqnfile = htm,  
               data = frequencies,  
               restrictions = list("dn=do", "g=.50"))
```

*# (B) hard-coding of constraints in the EQN file:*


```
htm_constr <- "  
target hit d  
target hit (1-d)*.50  
target miss (1-d)*.50  
  
lure cr d  
lure fa (1-d)*.50  
lure cr (1-d)*.50  
"  
fit <- traitMPT(eqnfile = htm_constr,  
               data = frequencies)
```

# Convergence


- Convergence check
  - Posterior/ MCMC samples should look unsystematic (like a hairy caterpillar)
  - For more options, see: `?plot.traitMPT`

```
plot(fit, parameter = "mean", type = "default")
```


Trace of mean[dn]




Density of mean[dn]



Trace of mean[g]




Density of mean[g]




# Convergence

## ■ Interpreting MCMC plots

**Good convergence: 'hairy caterpillar'**




**Bad convergence: slow movement**




# Convergence

- Gelman-Rubin statistic ( $\hat{R}$ )
  - Also known as: “potential scale reduction factor” or “R hat”
  - Similar to ANOVA: Compares between-chain and within-chain variances (large differences between these variances indicate nonconvergence)
  - Statistic should be close to 1 (standard criterion:  $\hat{R} < 1.05$ )

**Good convergence: all chains similar**



**Bad convergence: chains differ**



# Convergence

- Gelman-Rubin statistic ( $\hat{R}$ )
  - Columns Rhat and R\_95% in the summary output

```
summary(fit)
```

```
## Call:
## traitMPT(eqnfile = htm, data = frequencies, restrictions = list("dn=do"),
##         ppp = 1000)
##
## Group-level medians of MPT parameters (probability scale):
##           Mean    SD  2.5%  50% 97.5% Time-series SE n.eff  Rhat R_95%
## mean_dn 0.653 0.046 0.560 0.653 0.742          0.002   421 1.002 1.006
## mean_g  0.497 0.032 0.434 0.497 0.560          0.001   979 1.011 1.039
##
## Mean/Median of latent-trait values (probit-scale) across individuals:
##           Mean    SD  2.5%  50% 97.5% Time-series SE n.eff  Rhat R_95%
## latent_mu_dn 0.396 0.127 0.151 0.393 0.648          0.006   419 1.002 1.006
## latent_mu_g -0.007 0.080 -0.167 -0.008 0.151          0.003   979 1.011 1.039
##
## Standard deviation of latent-trait values (probit scale) across individuals:
##           Mean    SD  2.5%  50% 97.5% Time-series SE n.eff  Rhat R_95%
## latent_sigma_dn 0.857 0.110 0.672 0.847 1.106          0.002  3401 1.000 1.001
## latent_sigma_g  0.424 0.069 0.303 0.419 0.573          0.001  4517 1.001 1.005
##
## Correlations of latent-trait values on probit scale:
##           Mean    SD  2.5%  50% 97.5% Time-series SE n.eff  Rhat R_95%
## rho[dn,g] 0.045 0.206 -0.364 0.047 0.436          0.005  1992 1.004 1.016
##
## Correlations (posterior mean estimates) in matrix form:
```



## Options for MCMC Sampling

- If the model has not converged, it must be fitted with more conservative settings:

```
fit <- traitMPT(  
  eqnfile = htm, data = frequencies,  
  restrictions = list("dn=do"),  
  
  n.adapt = 5000, # longer adaption of JAGS increases efficiency of sampling  
  n.burnin = 5000, # longer burnin avoids issues due to bad starting values  
  n.iter = 30000, # drawing more MCMC samples leads to higher precision  
  n.thin = 10,    # ommitting every 10th sample reduces memory load  
  n.chains = 4)  # more MCMC chains increase precision
```

```
## MCMC sampling started at 2022-05-09 13:20:32  
## Calling 4 simulations using the parallel method...  
## Following the progress of chain 1 (the program will wait for all chains  
## to finish before continuing):  
## Welcome to JAGS 4.3.1 on Mon May 9 13:20:36 2022  
## JAGS is free software and comes with ABSOLUTELY NO WARRANTY  
## Loading module: basemod: ok  
## Loading module: bugs: ok  
## . Loading module: dic: ok  
## . Loading module: glm: ok  
## . . Reading data file data.txt  
## . Compiling data graph  
## Resolving undeclared variables
```

# Extend MCMC Sampling

- If the MCMC samples are OK but higher precision is needed:
  - Extend sampling and add new MCMC samples to the fitted JAGS object

```
fit2 <- extendMPT(fit,           # fitted MPT model
                 n.adapt = 2000, # JAGS need to restart and adapt again
                 n.burnin = 0,   # burnin not needed if previous samples are OK
                 n.iter = 10000) # how many additional iterations?
```


```
## Calling 4 simulations using the parallel method...
## Following the progress of chain 1 (the program will wait for all chains
## to finish before continuing):
## Welcome to JAGS 4.3.1 on Mon May  9 13:21:09 2022
## JAGS is free software and comes with ABSOLUTELY NO WARRANTY
## Loading module: basemod: ok
## Loading module: bugs: ok
## . Loading module: dic: ok
## . Loading module: glm: ok
## . . Reading data file data.txt
## . Compiling data graph
##   Resolving undeclared variables
##   Allocating nodes
##   Initializing
##   Reading data back into data table
## Compiling model graph
##   Resolving undeclared variables
##   Allocating nodes
## Graph information:
##   Observed stochastic nodes: 100
##   Unobserved stochastic nodes: 55
##   Total graph size: 1094
```

# Convergence for Second Fit


- Check convergence again:

```
plot(fit2, parameter = "mean", type = "default")
```


Trace of mean[dn]




Density of mean[dn]



Trace of mean[g]



Density of mean[g]



# Parameter Estimates

- Summary statistics for posterior distribution:
  - Posterior mean and median (50% quantile)
  - Posterior standard deviation (SD, similar to standard error)
  - Bayesian credibility interval (2.5% and 97.5% quantiles)

```
summary(fit2)
```


```
## Call:
## [[1]]
## traitMPT(eqnfile = htm, data = frequencies, restrictions = list("dn=do"),
##   n.iter = 30000, n.adapt = 5000, n.burnin = 5000, n.thin = 10,
##   n.chains = 4)
##
## [[2]]
## extendMPT(fittedModel = fit, n.iter = 10000, n.adapt = 2000,
##   n.burnin = 0)
##
##
## Group-level medians of MPT parameters (probability scale):
##      Mean    SD  2.5%  50% 97.5% Time-series SE n.eff  Rhat R_95%
## mean_dn 0.655 0.045 0.565 0.656 0.740      0.001  1125 1.006 1.016
## mean_g  0.500 0.032 0.437 0.500 0.562      0.001  1778 1.004 1.011
##
## Mean/Median of latent-trait values (probit-scale) across individuals:
##      Mean    SD  2.5%  50% 97.5% Time-series SE n.eff  Rhat R_95%
## latent_mu_dn 0.402 0.124 0.163 0.402 0.643      0.004  1122 1.006 1.016
## latent_mu_g -0.001 0.080 -0.158 -0.001 0.156      0.002  1778 1.004 1.011
##
## Standard deviation of latent-trait values (probit scale) across individuals:
##      Mean    SD  2.5%  50% 97.5% Time-series SE n.eff  Rhat R_95%
## latent_sigma_dn 0.855 0.110 0.667 0.845 1.098      0.002  4822 1.001 1.003
## latent_sigma_g 0.425 0.068 0.306 0.420 0.572      0.001  7352 1.000 1.001
```

# Plot Parameter Estimates

- Estimates for group-level parameters
  - Overall mean  $\mu$ : Posterior mean and Bayesian credibility interval
  - Individual parameters  $\theta_i$ : Posterior mean

```
plotParam(fit)
```

**roup-level means + 95% CI (red) and individual means**




# Plot Parameter Estimates

- Plot parameter profiles
  - E.g., assess test-retest reliability of a parameter (Michalkiewicz & Erdfelder, 2016)
- For more options, see: `?plotParam`

```
plotParam(fit, addLines = TRUE, select = c("dn", "g"))
```

**roup-level means + 95% CI (red) and individual means**





# Compare Prior and Posterior

## How much did we learn about the parameters?

- Graphical assessment: Plot prior (blue) and posterior (black) densities

```
plotPriorPost(fit)
```

```
## Press <Enter> to show the next plot.
```




```
## Press <Enter> to show the next plot.
```

# Group-Level Distribution

## Distribution of individual estimates

- Histogram: Distribution of  $\theta$  estimates (posterior mean per person)
- Red density: Estimated group-level distribution

```
plotDistribution(fit) # graphical test
```





# Model Fit: Predicted vs. Observed Data

## Graphical test of model fit


- Plot means of observed frequencies
- Compare against posterior-predicted frequencies (boxplots)

```
colMeans(frequencies)  # observed group means that are tested
```

```
##      cr      fa      hit      miss  
## 40.76  9.24 40.48  9.52
```

```
plotFit(fit)           # graphical test
```

Observed (red) and predicted (boxplot) mean frequencies



# Model Fit: Predicted vs. Observed Data


## ■ Model fit for *covariance* of observed frequencies

```
cov(frequencies)           # observed covariance matrix that is tested
```

```
##           cr          fa          hit          miss
## cr      46.88000 -46.88000  23.26041 -23.26041
## fa     -46.88000  46.88000 -23.26041  23.26041
## hit     23.26041 -23.26041  60.25469 -60.25469
## miss   -23.26041  23.26041 -60.25469  60.25469
```

```
plotFit(fit, stat = "cov") # graphical test
```

**Observed (red) and predicted (gray) covariances**



# Model Fit: Predicted vs. Observed Data

## Testing model fit

- The statistics T1 and T2 quantify the discrepancy between observed and expected means/covariances (similar to Pearson's  $\chi^2$ )
- Posterior predictive p-values
  - Values around .50 indicate good model fit
  - Values close to 0 (or close to 1) indicate misfit
  - In contrast to frequentist  $p$ -values, PPP values are *not* uniformly distributed when generating data from the correct model

```
PPP(fit, M = 2000, nCPU = 4)
```

```
## ## Mean structure (T1):
## Observed = 0.0349425 ; Predicted = 0.03475574 ; p-value = 0.504
##
## ## Covariance structure (T2):
## Observed = 6.584864 ; Predicted = 7.173334 ; p-value = 0.5195
##
## ## Individual fit (T1):
##      1      2      3      4      5      6      7      8      9     10     11     12     13
## 0.514 0.512 0.281 0.426 0.429 0.274 0.508 0.531 0.516 0.438 0.525 0.490 0.498
##     14     15     16     17     18     19     20     21     22     23     24     25     26
## 0.529 0.490 0.411 0.528 0.542 0.506 0.551 0.428 0.541 0.550 0.483 0.472 0.438
##     27     28     29     30     31     32     33     34     35     36     37     38     39
## 0.354 0.546 0.498 0.510 0.594 0.501 0.514 0.559 0.517 0.523 0.526 0.562 0.476
##     40     41     42     43     44     45     46     47     48     49     50
## 0.602 0.520 0.528 0.537 0.431 0.526 0.424 0.292 0.401 0.563 0.244
```

## Modeling with TreeBUGS is simple

- 1 Define model and clean data
- 2 Draw MCMC samples
- 3 Check convergence
- 4 Check model fit
- 5 Interpret/plot parameters

Note that these are the usual steps in any Bayesian analysis. . .

## Advanced Modeling

## Assessing within-subject differences in parameters

- 1 Data: Additional columns for separate within-subject conditions
- 2 Model: Write EQN file for within-subject design
- 3 MCMC sampling (as usual)
- 4 Comparison: Compute differences of parameters (transformed parameters)

### (1) Data structure for within-subject design

```
freq_within <- read.csv("2htm_within.csv")  
head(freq_within, 3)
```

##		high_cr	high_fa	high_hit	high_miss	low_cr	low_fa	low_hit	low_miss
##	1	40	10	28	22	38	12	32	18
##	2	38	12	36	14	45	5	37	13
##	3	46	4	44	6	39	11	30	20

# Within-Subject Comparisons

## (2) Model: Function for writing within-subject EQN files

- TreeBUGS has a function extends a standard MPT model to multiple within-subject conditions
- Essentially, model equations are copied and each parameter gets a new label (e.g., d\_condition1)

```
# create EQN file for within-subject manipulations
withinSubjectEQN(htm_d,
                 labels = c("high", "low"), # factor labels
                 constant=c("g"))          # constant parameters
```

	Tree	Category	Equation
## 1	high_target	high_hit	d_high
## 2	high_target	high_hit	$(1-d\_high)*g$
## 3	high_target	high_miss	$(1-d\_high)*(1-g)$
## 4	high_lure	high_cr	d_high
## 5	high_lure	high_fa	$(1-d\_high)*g$
## 6	high_lure	high_cr	$(1-d\_high)*(1-g)$
## 7	low_target	low_hit	d_low
## 8	low_target	low_hit	$(1-d\_low)*g$
## 9	low_target	low_miss	$(1-d\_low)*(1-g)$
## 10	low_lure	low_cr	d_low
## 11	low_lure	low_fa	$(1-d\_low)*g$
## 12	low_lure	low_cr	$(1-d\_low)*(1-g)$

## (4) Transformed parameters

- Often the interest is in the difference of a parameter across conditions
  - Example: Difference in memory strength  $\Delta_d = d_{\text{high}} - d_{\text{low}}$
- Based on the MCMC samples, we can simply compute any function of interest
  - We get a new set of posterior samples that can be summarized as usual

```
# fit to all conditions:
fit_within <- traitMPT("2htm_within.eqn", "2htm_within.csv")

# compute difference in d:
diff_d <- transformedParameters(
  fit_within,
  transformedParameters = list("diff_d = d_high - d_low"),
  level = "group")
summary(diff_d)$statistics
```

##	Mean	SD	Naive SE	Time-series SE
##	0.3186321511	0.0347032382	0.0003339321	0.0010913773



# Between-Subject Comparisons

## Comparisons in between-subject designs

- 1) Fit MPT model to each condition separately
  - Separate group-level parameters  $\mu$  and  $\Sigma$  per group
- 2) Compute differences in group-level parameters across conditions/models

```
fit1 <- traitMPT(htm_d, "2htm.csv")
fit2 <- traitMPT(htm_d, "2htm_group2.csv")

diff_between <- betweenSubjectMPT(
  fit1, fit2,                # fitted MPT models
  par1 = "d",               # parameter to test
  stat = c("x-y", "x>y"))  # transformed parameters

diff_between
```

##	Mean	SD	2.5%	50%	97.5%	Time-series	SE	n.eff	Rhat	R_95%
## d.m1-d.m2	0.301	0.07	0.163	0.303	0.438		0.004	322	1.037	1.115
## d.m1>d.m2	1.000	0.00	1.000	1.000	1.000		NaN	0	NaN	NaN

## Regression of MPT parameters on covariates

- Example: Predict memory performance  $d$  as a function of age
- Statistically, this requires an regression extension to the model
- The latent probit values  $\theta'_i$  are predicted by a design matrix  $X$ :

$$\theta'_i = \mu + X_i\beta + \delta_i$$

## Implementation in TreeBUGS

- Requires only two new arguments to provide the data (age of persons) and the regression structure (predict parameter  $D_n$  by age)

```
fit_regression <- traitMPT(htm_d, data = "2htm.csv",  
                           covData = "covariates.csv",  
                           predStructure = list("d ; continuous"))
```

```
round(fit_regression$summary$group$slope[,-6], 2)
```

##	Mean	SD	2.5%	50%	97.5%	n.eff	Rhat	R_95%
## slope_d_continuous	0.19	0.05	0.09	0.19	0.29	511	1	1
## slope_std_d_continuous	0.49	0.10	0.27	0.50	0.66	610	1	1


# Bayes Factor for Covariate

## Compute a Bayes factor

- H0: Slope parameter  $\beta = 0$
- H1: Slope parameter  $\beta \sim \text{Cauchy}(0, r)$  (with scale parameter  $r$ )
- Method: Savage-Dickey density ratio (Wagenmakers, 2010)
  - Bayes factor H1 vs. H0: prior divided by posterior density (at  $\beta = 0$ )
  - Only works for simple regression with 1 predictor (Heck 2019)

```
BayesFactorSlope(fit_regression,  
                  parameter = "slope_d_continuous",  
                  direction = ">",      # H1: positive slope parameter  
                  plot = TRUE)         # plot Savage-Dickey density ratio
```

Bayes factor B\_10=284.609 (prior red; posterior blue)



```
##                               BF_0>    BF_>0  
## slope_d_continuous 0.003513589 284.6093
```

# Between-Subject Comparisons: Similar to ANOVA

## Between-subject designs: Assumptions about the covariance matrix $\Sigma$

- A) Separate covariance matrix per condition:  $\Sigma_1, \Sigma_2, \dots$ 
  - See previous slides: `betweenSubjectMPT(fit1, fit2)`
- B) Identical covariance matrix  $\Sigma$  across conditions
  - Similar to ANOVA: “pooled variance” (Rouder & Morey; 2012)
  - Manipulation only affects the mean parameters  $\mu$

```
# fit all between-conditions jointly:
fit_between <- traitMPT(
  htm_d, "2htm.csv",
  covData = "covariates.csv",
  predStructure = list("d ; discrete"), # discrete predictor
  predType = c("c", "f")) # "c" =continuous; "f"=fixed-effects
```

```
# get estimates for the group-specific MPT parameters
gmeans <- getGroupMeans(fit_between)
round(gmeans, 2)
```

```
##               Mean   SD 2.5%  50% 97.5% p(one-sided vs. overall)
## d_discrete[group_a] 0.63 0.07 0.49 0.63 0.75                      0.32
## d_discrete[group_b] 0.67 0.07 0.54 0.68 0.79                      0.32
```

## Sensitivity/robustness analysis

## Define different priors

- Prior distributions in the latent-trait MPT necessary for:
  - Latent (probit-) mean  $\mu$
  - Latent (probit-) covariance matrix  $\Sigma$ : scaled inverse Wishart with
    - Prior matrix  $V$
    - Degrees of freedom  $df$
    - Scaling parameter  $\xi$
- Example: We assume that guessing probabilities are around 50%

```
fit <- traitMPT(eqnfile="htm.txt", data="responses.csv",
               restrictions=list("dn=do"),

               mu = c(dn = "dnorm(0,1)",      # default prior
                      g = "dnorm(0,5)"),      # prior focused around 50% guessing
               xi = "dunif(0,2)",              # less dispersion of MPT parameters
               V = diag(2),                    # default
               df = 2 + 1)                     # default
```


# Understanding Priors

## What do the priors actually mean?


- Draw samples from the prior
- Plot mean/SD of MPT parameters

```
plotPrior(prior =  
  list(mu = c(dn = "dnorm(0,1)", # default prior  
             g = "dnorm(0,5)", # prior focused around 50%  
             xi="dunif(0,2)", # smaller scale for group SD  
             V= diag(2), df = 3)) # default Wishart prior
```


Prior on group mean:  $\mu = \text{dnorm}(0,1)$




Prior on group mean:  $\mu = \text{dnorm}(0,5)$



Prior on group SD:  $\xi = \text{dunif}(0,2)$



Prior on group SD:  $\xi = \text{dunif}(0,2)$




# Prior Predictive Sampling

## Prior predictive distribution

- 1 Draw samples from the prior
- 2 Draw new data (response frequencies)
- 3 Assess predicted data (e.g., plots or descriptive statistics)


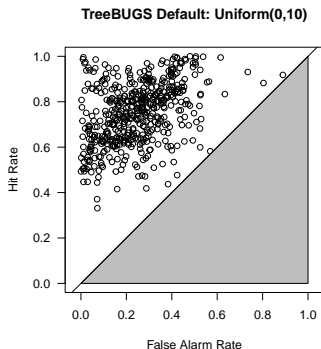
```
pp <- priorPredictive(prior = list(mu = "dnorm(0,1)", xi="dunif(0,10)",  
                                   V=diag(2), df=2+1),  
                      eqnfile = htm, restrictions = list("dn=do"),  
                      numItems = c(target = 50, lure = 50),  
                      N = 50, M = 500)      # number of participants/samples  
  
# compute and plot predicted values for the average hit/FA rates (in ROC space)  
mean_freq <- sapply(pp, colMeans)  
par(mar=c(4,5,.1, .1))  
plot(mean_freq["fa",,]/50, mean_freq["hit",,]/50, asp = 1, xlim = 0:1, ylim = 0:1)  
polygon(c(0,1,1), c(0,0,1), col = "gray")  
abline(0, 1)
```





## Excursion: Different default priors for the scale parameter $\xi$

- 1 TreeBUGS (Heck et al., 2018):  $\xi \sim \text{Uniform}(0, 10)$
- 2 Klauer (2010) and Matzke et al. (2015):  $\xi \sim \text{Uniform}(0, 100)$




# Posterior Predictive

## Posterior Predictive Distribution

- What data does the fitted model predict?
- Use posterior samples of the *parameters* to draw new samples of the *data* (i.e., predicted response frequencies)
  - Note: These are the basis of posterior-predictive checks (T1 and T2 statistics)

```
postpred <- posteriorPredictive(fit, M = 100, nCPU = 4)
```



## Sensitivity and robustness analysis

- Assessing the impact of priors, estimate necessary sample size for specific analysis etc.
  - 1 Generate data from (correct or wrong) model
  - 2 Fit model
  - 3 If necessary: Replicate with a for-loop

```
# standard, fixed-effects MPT (generate data for one person)
sim <- genMPT(theta = c(d = .7, g = .5),
              numItems = c(target = 50, lure = 50),
              eqnfile = htm_d)

# hierarchical MPT (generate a complete table of frequencies)
sim2 <- genTraitMPT(N = 100, eqnfile = htm_d,
                   mean = c(d = .7, g = .5),
                   sigma = c(d = .4, g = .2),
                   rho = diag(2))
```

## Appendix

# Appendix: Testing for Heterogeneity

## ■ Test by Smith and Batchelder (2008)

A Test person heterogeneity assuming items homogeneity ( $\chi^2$ )

B Test person heterogeneity under item heterogeneity (permutation bootstrap)

```
# A) chi^2 test
test <- testHetChi(freq = frequencies,
                  tree = c(cr="lure", fa="lure",
                          hit="target", miss="target"))
data.frame(test)
```

```
##      chisq df      prob
## 1 688.0345 98 2.099801e-89
```


```
# B) requires data in long format (variables: person / item / response)
# testHetPerm(data, tree, source = "person")
```

## Appendix: Convergence

### Autocorrelation function

- How strongly are the MCMC samples correlated between iteration  $t$  and iteration  $t + \text{Lag}$ ?
- Ideally, these curves should rapidly decrease towards zero.



```
plot(fit, parameter = "mean", type = "acf")
```



## Appendix: Convergence

- Plot evolution of Gelman-Rubin statistic
  - Also known as: “potential scale reduction factor” or “R hat”
  - Similar to ANOVA: Compares between-chain and within-chain variances (large differences between these variances indicate nonconvergence)
  - Statistic should be close to 1

```
plot(fit, parameter = "mean", type = "gelman")
```



### **MCMC samplers available in TreeBUGS**

- Fixed-effects (“standard”) MPT: Gibbs sampler in C++
- Beta-MPT: JAGS and C++
- Latent-trait with extensions: JAGS
  - Combination of random- and fixed-effects parameters
  - Continuous and categorical covariates
  - Group-level structure: Independent normal distributions



# References I

- Heck, Daniel W. 2019. "A Caveat on the Savage-Dickey Density Ratio: The Case of Computing Bayes Factors for Regression Parameters." *British Journal of Mathematical and Statistical Psychology* 72: 316–33.  
<https://doi.org/10.1111/bmsp.12150>.
- Heck, Daniel W, Nina R. Arnold, and Denis Arnold. 2018. "TreeBUGS: An R Package for Hierarchical Multinomial-Processing-Tree Modeling." *Behavior Research Methods* 50: 264–84. <https://doi.org/10.3758/s13428-017-0869-7>.
- Klauer, K. C. 2010. "Hierarchical Multinomial Processing Tree Models: A Latent-Trait Approach." *Psychometrika* 75: 70–98.  
<https://doi.org/10.1007/s11336-009-9141-0>.
- Matzke, Dora, Conor V. Dolan, William H. Batchelder, and Eric-Jan Wagenmakers. 2015. "Bayesian Estimation of Multinomial Processing Tree Models with Heterogeneity in Participants and Items." *Psychometrika* 80: 205–35. <https://doi.org/10.1007/s11336-013-9374-9>.
- Smith, J. B., and W. H. Batchelder. 2008. "Assessing Individual Differences in Categorical Data." *Psychonomic Bulletin & Review* 15: 713–31.  
<https://doi.org/10.3758/PBR.15.4.713>.